

## Introduction to noSQL databases Part-2

### Using cassandra on cyverse

March 11, 2024

Shashank




RESEARCH, INNOVATION & IMPACT  
Data Science Institute

Cassandra is based on a distributed system architecture.

In its simplest form, Cassandra can be installed on a single machine or container. A single Cassandra instance is called a node.

Cassandra supports horizontal scalability achieved by adding more than one node as a part of a Cassandra cluster.



Database

MongoDB

Apache Cassandra

Use Cases

- Search use cases
- E-commerce websites

- "Always available" type of services (e.g., Netflix, Spotify)

Features

- Read with indexes
- Consistency
- Primary-Secondary architecture

- Fast writes - capture all data
- Availability & scalability



RESEARCH, INNOVATION & IMPACT

Data Science Institute



Examples on cyverse

